

An Implementation of Scheduler Activations on the NetBSD Operating System

Nathan J. Williams
Wasabi Systems, Inc.
nathanw@wasabisystems.com

Abstract

This paper presents the design and implementation of a two-level thread scheduling system on NetBSD. This system provides a foundation for efficient and flexible threads on both uniprocessor and multiprocessor machines. The work is based on the scheduler activations kernel interface proposed by Anderson et al. [1] for user-level control of parallelism in the presence of multiprocessing and multiprocessing.

Introduction

Thread programming has become a popular and important part of application development. Some programs want to improve performance by exploiting concurrency; others find threads a natural way to decompose their application structure. However, the two major types of thread implementation available - user threads and kernel threads - both have significant drawbacks in overhead and concurrency that limit the performance of applications.

A potential solution to this problem exists in the form of a hybrid, two-level thread system known as *scheduler activations* that divides the work between the kernel and the user levels. Such a system has the potential to achieve the high performance of user threads while retaining the concurrency of kernel threads.

The purpose of this paper is to describe the design and implementation of such a two-level system and associated thread library in NetBSD, show that the speed of thread operations is competitive with user thread implementations, and demonstrate that it can be implemented without hurting the performance of unrelated parts of the system.

First, as motivation, Section 2 describes traditional thread implementations. Scheduler activations are explained in the context of two-level thread implementations, and then described in detail in Section 3, giving the interface that the scheduler activations system presents to programs, and Section 5 follows with details about the kernel implementation behind the interface. The thread library built on this interface is described in Section 6, and the performance of the system and the thread library is examined and compared to other libraries in Section 7. Finally, Section 8 concludes and considers directions for future work.

Thread Systems

Historically, there have been two major types of thread implementations on Unix-like systems, with the essential differences being the participation of the kernel in the thread management. Each type has significant drawbacks, and much work has gone into finding a compromise or third way.

The first type of thread system is implemented purely at user-level. In this system, all thread operations manipulate state that is private to the process, and the kernel is unaware of the presence of the threads. This type of thread system is often known as “N:1” thread system because the thread implementation maps all of the N application threads onto a single kernel resource.

Examples of this type of thread system include the GNU PTH thread library [6], FSU Pthreads, PTL2 [7], Provenzano’s “MIT pthreads” library [8], and the original DCE threads package that formed so much of the basis for the POSIX thread effort [4]. Some large software packages contain their own thread system, especially those that were originally written to support platforms without native thread support (such as the “green threads” in Sun’s original Java implementation). All of the “BSDs” currently have one of these user-level packages as their primary thread system.

User-level threads can be implemented without kernel support, which is useful for platforms without native thread support, or applications where only a particular subset of thread operations is needed. The GNU PTH library, for example, does not support preemptive time-slicing among threads, and is simpler because of it. Thread creation, synchronization, and context switching can all be implemented with a cost comparable to an ordinary function call.

However, operations that conceptually block a single thread (blocking system calls such as read(), or page faults) instead block the entire process, since the kernel is oblivious to the presence of threads. This makes it difficult to use user-level threads to exploit concurrency or provide good interactive response. Many user-level thread packages partially work around this problem by intercepting system calls made by the application and replacing them with non-blocking variants and a call to the thread scheduler. These workarounds are not entirely effective and add complexity to the system.

Additionally, a purely user-level thread package can not make use of multiple CPUs in a system. The kernel is only aware of one entity that can be scheduled to run - the process - and hence only allocates a single processor. As a result, user-level thread packages are unsuitable for applications that are natural fits for shared-memory multiprocessors, such as large numerical simulations.

At the other end of the thread implementation spectrum, the operating system kernel is aware of the threaded nature of the application and the existence of each application thread. This model is known as the “1:1” model, since there is a direct correspondence between user threads and kernel resources. The kernel is responsible for most thread management tasks: creation, scheduling, synchronization, and disposal. These kernel entities share many of the resources traditionally associated with a process, such as address space and file descriptors, but each have their own running state or saved context.

This approach provides the kernel with awareness of the concurrency that exists within an application. Several benefits are realized over the user-thread model: one thread blocking does not impede the progress of another, and multiprocessor parallelism can be exploited. But there are problems here as well. One is that the overhead of thread operations is high: since they are managed by the kernel, operations must be performed by requesting services of the kernel (usually via system call), which is a relatively slow operation. Also, each thread consumes kernel memory, which is usually more scarce than user process memory. Thus, while kernel threads provide better concurrency than user threads, they are more expensive in time and space. They are relatively easy to implement, given operating system support for kernel execution entities that share resources (such as the clone() system call under Linux, the spawn() system call under IRIX, or the _wup_create() system call in Solaris). Many operating systems, including Linux, IRIX, and Windows NT, use this model of thread system.

Since there are advantages and disadvantages of both the N:1 and 1:1 thread implementation models, it is natural to attempt to combine them to achieve a balance of the costs and benefits of each. These hybrids are collectively known as “N:M” systems, since they map some number N of application threads onto a (usually smaller) number M of kernel entities. They are also known as “two-level” thread systems, since there are two parties, the kernel and the user parts of the thread system, involved in thread operations and scheduling. There are quite a variety of different implementations of N:M thread systems, with different performance characteristics. N:M thread systems are more complicated than either of the other models, and can be more difficult to develop, debug, and use effectively. Both AIX and Solaris use N:M thread systems by default. 2

In a N:M thread system, a key question is how to manage the mapping of user threads to kernel entities. One possibility is to associate groups of threads with single kernel entities; this permits concurrency across groups but not within groups, reaching a balance between the concurrency of N:1 and 1:1 systems.

The *scheduler activations* model put forward by Anderson et al. is a way of managing the N:M mapping while maintaining as much concurrency as a 1:1 thread system. In this model, the kernel provides the application with the abstraction of virtual processors: a guarantee to run a certain number of application threads simultaneously on CPUs. Kernel events that would affect the number of running threads are communicated directly to the application in a way that maintains the number of virtual processors. The message to the application informs it of the state that has changed and the context of the user threads that were involved, and lets the user-level scheduler decide how to proceed with the resources available.

This system has several advantages: like other M:N systems, kernel resource usages is kept simple in comparison to the number of user-level threads; voluntary thread switching is cheap, similar to user-level threads, and like 1:1 systems, an application’s concurrency is fully maintained. Scheduler activations have been implemented for research purposes in Taos [1], Mach 3.0 [2], and BSD/OS [9], and adopted commercially in Digital Unix [5] (now Compaq Tru64 Unix).

The scheduler activations system shares with other M:N systems all of the problems of increased complexity over 1:1 systems. Additionally, there is concern that the problems addressed by scheduler activations are not important problems in the space of threaded applications. For example, making thread context switches cheap is of little value if thread-to-thread switching is infrequent, or if thread switching occurs as a side effect of heavyweight I/O operations.

Implementing scheduler activations for NetBSD is attractive for two major reasons. First, NetBSD needs a native thread system which is preemptive and has the ability to exploit multiprocessor computer systems. Second, this work makes a scheduler activations interface and implementation available in an open-source operating system for continued research into the utility and viability of this intuitively appealing model.

Scheduler Activations

As described by Anderson et al., the scheduler activations kernel provides the application with a set of virtual processors, and then the application has complete control over what threads to run on each of the virtual processors. The number of virtual processors in the set is controlled by the kernel, in response to the competing demands of different processes in the system. For example, an application may express to the kernel that it has enough work to keep four processors busy, while a single-threaded application is also trying to run; the kernel could allocate three processors to the set of virtual processors for the first application, and give the fourth processor to the single-threaded program.

In order for the application to be able to consistently use these virtual processors, it must know when threads have blocked, stopped, or restarted. For user-level operations that cause threads to block, such as sleeping for a mutex or waiting on a condition variable, the thread that is blocking hands control to the thread library, which can schedule another thread to run in the usual manner. However, kernel-level events can also block threads: the read() and select() system calls, for example, or taking a fault on a page of memory that is on disk. When such an event occurs, the number of processors executing application code decreases. The scheduler activations kernel needs to tell the application what has happened and give it another virtual processor. The mechanism that does this is known as an *upcall*.

To perform an upcall, the kernel allocates a new virtual processor for the application and begins running a piece of application code in this new execution context. The application code is known as the *upcall handler*, and it is invoked similarly to a traditional signal handler. The upcall handler is passed information that identifies the virtual processor that stopped running and the reason that it stopped. The upcall handler can then perform any user-level thread bookkeeping and then switch to a runnable thread from the application’s thread queue.

Eventually the thread that blocked in the kernel will unblock and be ready to return to the application. If the thread were to directly return, it would violate two constraints of scheduler activations: the number of virtual processors would increase, and the application would be unaware that the state of the thread had changed. Therefore, this event is also communicated with an upcall. In order to maintain the number of virtual processors, the thread currently executing on one of the application’s processors is preempted. The upcall is then run in the context of that virtual processor, carrying notifications of both the first thread returning from the kernel and the second thread being preempted. The upcall handler, knowing all of this, can then decide which thread to run next, based on its own processing needs. Typically, this would choose the highest priority thread, which should be one of the two involved in the upcall notification.

There are other scheduling-related events that are communicated by upcall. A change in the size of the virtual processor set must be communicated to the application so that it can either reschedule the thread running on a removed processor, or schedule code to run on the new processor. Traditional POSIX signals, which can only normally cause a thread to flow change in the application, are communicated by upcall. Additionally, a mechanism is provided for an application to invoke an upcall on another processor, in order to bring that processor back under control of the thread engine (in case thread engine code running on processor 1 decides that a different, higher-priority thread should start running on processor 2).

Kernel Interface

The application interface to the scheduler activations system consists of system calls. First, the `sa_register()` call tells the kernel what entry point to use for a scheduler activations upcall, much like registering a signal handler. Next, `sa_setconcurrency()` informs the kernel of the level of concurrency available in the application, and thus the maximum number of processors that may be profitably allocated to it. The `sa_enable()` call starts the system by invoking an upcall on the current processor. While the application is running, the `sa_yield()` and `sa_suspend()` calls allow an application to manage itself by giving up processors and interrupting other processors in the application with an upcall.

Upcalls

Upcalls are the interface used by the scheduler activations system in the kernel to inform an application of a scheduling-related event. When an upcall that makes use of scheduler activations registers a procedure to handle the upcall, much like registering a signal handler. When an event occurs, the kernel will take a processor allocated to the application (possibly preempting another part of the application), switch to user level, and call the registered procedure.

The signature of an upcall is:

```
void sa_upcall(int type,
               struct sa_t *sa_t,
               int events,
               int interrupted,
               void *arg);
```

The `type` argument indicates the event which triggered the upcall. The types and their meanings are described below.

The `sa_t` field points to an array of pointers to `struct sa_t` which describe the activations involved in this event. The first element of this array, `sa_t[0]`, points to the `sa_t` of the running activation. The next elements of the array (`sa_t[1]` through `sa_t[event+1]`) describe the activations that were preempted, and the thread currently executing on one of the application’s processors is preempted. The upcall is then run in the context of that virtual processor, carrying notifications of both the first thread returning from the kernel and the second thread being preempted. The upcall handler, knowing all of this, can then decide which thread to run next, based on its own processing needs. Typically, this would choose the highest priority thread, which should be one of the two involved in the upcall notification.

The `events` field is a bit mask of events that occurred. The `interrupted` field is a bit mask of events that occurred. The `arg` parameter points to a `siginfo_t` structure that describes the signal being delivered.

The `interrupted` field is a bit mask of events that occurred. The `arg` parameter points to a `siginfo_t` structure that describes the signal being delivered.

The `sa_t` field points to an array of pointers to `struct sa_t` which describe the activations involved in this event. The first element of this array, `sa_t[0]`, points to the `sa_t` of the running activation. The next elements of the array (`sa_t[1]` through `sa_t[event+1]`) describe the activations that were preempted, and the thread currently executing on one of the application’s processors is preempted. The upcall is then run in the context of that virtual processor, carrying notifications of both the first thread returning from the kernel and the second thread being preempted. The upcall handler, knowing all of this, can then decide which thread to run next, based on its own processing needs. Typically, this would choose the highest priority thread, which should be one of the two involved in the upcall notification.

The `events` field is a bit mask of events that occurred. The `interrupted` field is a bit mask of events that occurred. The `arg` parameter points to a `siginfo_t` structure that describes the signal being delivered.

The `interrupted` field is a bit mask of events that occurred. The `arg` parameter points to a `siginfo_t` structure that describes the signal being delivered.

The `sa_t` field points to an array of pointers to `struct sa_t` which describe the activations involved in this event. The first element of this array, `sa_t[0]`, points to the `sa_t` of the running activation. The next elements of the array (`sa_t[1]` through `sa_t[event+1]`) describe the activations that were preempted, and the thread currently executing on one of the application’s processors is preempted. The upcall is then run in the context of that virtual processor, carrying notifications of both the first thread returning from the kernel and the second thread being preempted. The upcall handler, knowing all of this, can then decide which thread to run next, based on its own processing needs. Typically, this would choose the highest priority thread, which should be one of the two involved in the upcall notification.

The `events` field is a bit mask of events that occurred. The `interrupted` field is a bit mask of events that occurred. The `arg` parameter points to a `siginfo_t` structure that describes the signal being delivered.

The `interrupted` field is a bit mask of events that occurred. The `arg` parameter points to a `siginfo_t` structure that describes the signal being delivered.

The `sa_t` field points to an array of pointers to `struct sa_t` which describe the activations involved in this event. The first element of this array, `sa_t[0]`, points to the `sa_t` of the running activation. The next elements of the array (`sa_t[1]` through `sa_t[event+1]`) describe the activations that were preempted, and the thread currently executing on one of the application’s processors is preempted. The upcall is then run in the context of that virtual processor, carrying notifications of both the first thread returning from the kernel and the second thread being preempted. The upcall handler, knowing all of this, can then decide which thread to run next, based on its own processing needs. Typically, this would choose the highest priority thread, which should be one of the two involved in the upcall notification.

The `events` field is a bit mask of events that occurred. The `interrupted` field is a bit mask of events that occurred. The `arg` parameter points to a `siginfo_t` structure that describes the signal being delivered.

The `interrupted` field is a bit mask of events that occurred. The `arg` parameter points to a `siginfo_t` structure that describes the signal being delivered.

The `sa_t` field points to an array of pointers to `struct sa_t` which describe the activations involved in this event. The first element of this array, `sa_t[0]`, points to the `sa_t` of the running activation. The next elements of the array (`sa_t[1]` through `sa_t[event+1]`) describe the activations that were preempted, and the thread currently executing on one of the application’s processors is preempted. The upcall is then run in the context of that virtual processor, carrying notifications of both the first thread returning from the kernel and the second thread being preempted. The upcall handler, knowing all of this, can then decide which thread to run next, based on its own processing needs. Typically, this would choose the highest priority thread, which should be one of the two involved in the upcall notification.

The `events` field is a bit mask of events that occurred. The `interrupted` field is a bit mask of events that occurred. The `arg` parameter points to a `siginfo_t` structure that describes the signal being delivered.

The `interrupted` field is a bit mask of events that occurred. The `arg` parameter points to a `siginfo_t` structure that describes the signal being delivered.

The `sa_t` field points to an array of pointers to `struct sa_t` which describe the activations involved in this event. The first element of this array, `sa_t[0]`, points to the `sa_t` of the running activation. The next elements of the array (`sa_t[1]` through `sa_t[event+1]`) describe the activations that were preempted, and the thread currently executing on one of the application’s processors is preempted. The upcall is then run in the context of that virtual processor, carrying notifications of both the first thread returning from the kernel and the second thread being preempted. The upcall handler, knowing all of this, can then decide which thread to run next, based on its own processing needs. Typically, this would choose the highest priority thread, which should be one of the two involved in the upcall notification.

The `events` field is a bit mask of events that occurred. The `interrupted` field is a bit mask of events that occurred. The `arg` parameter points to a `siginfo_t` structure that describes the signal being delivered.

The `interrupted` field is a bit mask of events that occurred. The `arg` parameter points to a `siginfo_t` structure that describes the signal being delivered.

The `sa_t` field points to an array of pointers to `struct sa_t` which describe the activations involved in this event. The first element of this array, `sa_t[0]`, points to the `sa_t` of the running activation. The next elements of the array (`sa_t[1]` through `sa_t[event+1]`) describe the activations that were preempted, and the thread currently executing on one of the application’s processors is preempted. The upcall is then run in the context of that virtual processor, carrying notifications of both the first thread returning from the kernel and the second thread being preempted. The upcall handler, knowing all of this, can then decide which thread to run next, based on its own processing needs. Typically, this would choose the highest priority thread, which should be one of the two involved in the upcall notification.

The `events` field is a bit mask of events that occurred. The `interrupted` field is a bit mask of events that occurred. The `arg` parameter points to a `siginfo_t` structure that describes the signal being delivered.

The `interrupted` field is a bit mask of events that occurred. The `arg` parameter points to a `siginfo_t` structure that describes the signal being delivered.

The `sa_t` field points to an array of pointers to `struct sa_t` which describe the activations involved in this event. The first element of this array, `sa_t[0]`, points to the `sa_t` of the running activation. The next elements of the array (`sa_t[1]` through `sa_t[event+1]`) describe the activations that were preempted, and the thread currently executing on one of the application’s processors is preempted. The upcall is then run in the context of that virtual processor, carrying notifications of both the first thread returning from the kernel and the second thread being preempted. The upcall handler, knowing all of this, can then decide which thread to run next, based on its own processing needs. Typically, this would choose the highest priority thread, which should be one of the two involved in the upcall notification.

The `events` field is a bit mask of events that occurred. The `interrupted` field is a bit mask of events that occurred. The `arg` parameter points to a `siginfo_t` structure that describes the signal being delivered.

The `interrupted` field is a bit mask of events that occurred. The `arg` parameter points to a `siginfo_t` structure that describes the signal being delivered.

The `sa_t` field points to an array of pointers to `struct sa_t` which describe the activations involved in this event. The first element of this array, `sa_t[0]`, points to the `sa_t` of the running activation. The next elements of the array (`sa_t[1]` through `sa_t[event+1]`) describe the activations that were preempted, and the thread currently executing on one of the application’s processors is preempted. The upcall is then run in the context of that virtual processor, carrying notifications of both the first thread returning from the kernel and the second thread being preempted. The upcall handler, knowing all of this, can then decide which thread to run next, based on its own processing needs. Typically, this would choose the highest priority thread, which should be one of the two involved in the upcall notification.

The `events` field is a bit mask of events that occurred. The `interrupted` field is a bit mask of events that occurred. The `arg` parameter points to a `siginfo_t` structure that describes the signal being delivered.

The `interrupted` field is a bit mask of events that occurred. The `arg` parameter points to a `siginfo_t` structure that describes the signal being delivered.

The `sa_t` field points to an array of pointers to `struct sa_t` which describe the activations involved in this event. The first element of this array, `sa_t[0]`, points to the `sa_t` of the running activation. The next elements of the array (`sa_t[1]` through `sa_t[event+1]`) describe the activations that were preempted, and the thread currently executing on one of the application’s processors is preempted. The upcall is then run in the context of that virtual processor, carrying notifications of both the first thread returning from the kernel and the second thread being preempted. The upcall handler, knowing all of this, can then decide which thread to run next, based on its own processing needs. Typically, this would choose the highest priority thread, which should be one of the two involved in the upcall notification.

The `events` field is a bit mask of events that occurred. The `interrupted` field is a bit mask of events that occurred. The `arg` parameter points to a `siginfo_t` structure that describes the signal being delivered.

The `interrupted` field is a bit mask of events that occurred. The `arg` parameter points to a `siginfo_t` structure that describes the signal being delivered.

The `sa_t` field points to an array of pointers to `struct sa_t` which describe the activations involved in this event. The first element of this array, `sa_t[0]`, points to the `sa_t` of the running activation. The next elements of the array (`sa_t[1]` through `sa_t[event+1]`) describe the activations that were preempted, and the thread currently executing on one of the application’s processors is preempted. The upcall is then run in the context of that virtual processor, carrying notifications of both the first thread returning from the kernel and the second thread being preempted. The upcall handler, knowing all of this, can then decide which thread to run next, based on its own processing needs. Typically, this would choose the highest priority thread, which should be one of the two involved in the upcall notification.

The `events` field is a bit mask of events that occurred. The `interrupted` field is a bit mask of events that occurred. The `arg` parameter points to a `siginfo_t` structure that describes the signal being delivered.

The `interrupted` field is a bit mask of events that occurred. The `arg` parameter points to a `siginfo_t` structure that describes the signal being delivered.

The `sa_t` field points to an array of pointers to `struct sa_t` which describe the activations involved in this event. The first element of this array, `sa_t[0]`, points to the `sa_t` of the running activation. The next elements of the array (`sa_t[1]` through `sa_t[event+1]`) describe the activations that were preempted, and the thread currently executing on one of the application’s processors is preempted. The upcall is then run in the context of that virtual processor, carrying notifications of both the first thread returning from the kernel and the second thread being preempted. The upcall handler, knowing all of this, can then decide which thread to run next, based on its own processing needs. Typically, this would choose the highest priority thread, which should be one of the two involved in the upcall notification.

The `events` field is a bit mask of events that occurred. The `interrupted` field is a bit mask of events that occurred. The `arg` parameter points to a `siginfo_t` structure that describes the signal being delivered.

The `interrupted` field is a bit mask of events that occurred. The `arg` parameter points to a `siginfo_t` structure that describes the signal being delivered.

The `sa_t` field points to an array of pointers to `struct sa_t` which describe the activations involved in this event. The first element of this array, `sa_t[0]`, points to the `sa_t` of the running activation. The next elements of the array (`sa_t[1]` through `sa_t[event+1]`) describe the activations that were preempted, and the thread currently executing on one of the application’s processors is preempted. The upcall is then run in the context of that virtual processor, carrying notifications of both the first thread returning from the kernel and the second thread being preempted. The upcall handler, knowing all of this, can then decide which thread to run next, based on its own processing needs. Typically, this would choose the highest priority thread, which should be one of the two involved in the upcall notification.

The `events` field is a bit mask of events that occurred. The `interrupted` field is a bit mask of events that occurred. The `arg` parameter points to a `siginfo_t` structure that describes the signal being delivered.

The `interrupted` field is a bit mask of events that occurred. The `arg` parameter points to a `siginfo_t` structure that describes the signal being delivered.

The `sa_t` field points to an array of pointers to `struct sa_t` which describe the activations involved in this event. The first element of this array, `sa_t[0]`, points to the `sa_t` of the running activation. The next elements of the array (`sa_t[1]` through `sa_t[event+1]`) describe the activations that were preempted, and the thread currently executing on one of the application’s processors is preempted. The upcall is then run in the context of that virtual processor, carrying notifications of both the first thread returning from the kernel and the second thread being preempted. The upcall handler, knowing all of this, can then decide which thread to run next, based on its own processing needs. Typically, this would choose the highest priority thread, which should be one of the two involved in the upcall notification.

The `events` field is a bit mask of events that occurred. The `interrupted` field is a bit mask of events that occurred. The `arg` parameter points to a `siginfo_t` structure that describes the signal being delivered.

The `interrupted` field is a bit mask of events that occurred. The `arg` parameter points to a `siginfo_t` structure that describes the signal being delivered.

The `sa_t` field points to an array of pointers to `struct sa_t` which describe the activations involved in this event. The first element of this array, `sa_t[0]`, points to the `sa_t` of the running activation. The next elements of the array (`sa_t[1]` through `sa_t[event+1]`) describe the activations that were preempted, and the thread currently executing on one of the application’s processors is preempted. The upcall is then run in the context of that virtual processor, carrying notifications of both the first thread returning from the kernel and the second thread being preempted. The upcall handler, knowing all of this, can then decide which thread to run next, based on its own processing needs. Typically, this would choose the highest priority thread, which should be one of the two involved in the upcall notification.

The `events` field is a bit mask of events that occurred. The `interrupted` field is a bit mask of events that occurred. The `arg` parameter points to a `siginfo_t` structure that describes the signal being delivered.

The `interrupted` field is a bit mask of events that occurred. The `arg` parameter points to a `siginfo_t` structure that describes the signal being delivered.

The `sa_t` field points to an array of pointers to `struct sa_t` which describe the activations involved in this event. The first element of this array, `sa_t[0]`, points to the `sa_t` of the running activation. The next elements of the array (`sa_t[1]` through `sa_t[event+1]`) describe the activations that were preempted, and the thread currently executing on one of the application’s processors is preempted. The upcall is then run in the context of that virtual processor, carrying notifications of both the first thread returning from the kernel and the second thread being preempted. The upcall handler, knowing all of this, can then decide which thread to run next, based on its own processing needs. Typically, this would choose the highest priority thread, which should be one of the two involved in the upcall notification.

The `events` field is a bit mask of events that occurred. The `interrupted` field is a bit mask of events that occurred. The `arg` parameter points to a `siginfo_t` structure that describes the signal being delivered.

The `interrupted` field is a bit mask of events that occurred. The `arg` parameter points to a `siginfo_t` structure that describes the signal being delivered.

The `sa_t` field points to an array of pointers to `struct sa_t` which describe the activations involved in this event. The first element of this array, `sa_t[0]`, points to the `sa_t` of the running activation. The next elements of the array (`sa_t[1]` through `sa_t[event+1]`) describe the activations that were preempted, and the thread currently executing on one of the application’s processors is preempted. The upcall is then run in the context of that virtual processor, carrying notifications of both the first thread returning from the kernel and the second thread being preempted. The upcall handler, knowing all of this, can then decide which thread to run next, based on its own processing needs. Typically, this would choose the highest priority thread, which should be one of the two involved in the upcall notification.

The `events` field is a bit mask of events that occurred. The `interrupted` field is a bit mask of events that occurred. The `arg` parameter points to a `siginfo_t` structure that describes the signal being delivered.

The `interrupted` field is a bit mask of events that occurred. The `arg` parameter points to a `siginfo_t` structure that describes the signal being delivered.

The `sa_t` field points to an array of pointers to `struct sa_t` which describe the activations involved in this event. The first element of this array, `sa_t[0]`, points to the `sa_t` of the running activation. The next elements of the array (`sa_t[1]` through `sa_t[event+1]`) describe the activations that were preempted, and the thread currently executing on one of the application’s processors is preempted. The upcall is then run in the context of that virtual processor, carrying notifications of both the first thread returning from the kernel and the second thread being preempted. The upcall handler, knowing all of this, can then decide which thread to run next, based on its own processing needs. Typically, this would choose the highest priority thread, which should be one of the two involved in the upcall notification.

The `events` field is a bit mask of events that occurred. The `interrupted` field is a bit mask of events that occurred. The `arg` parameter points to a `siginfo_t` structure that describes the signal being delivered.

The `interrupted` field is a bit mask of events that occurred. The `arg` parameter points to a `siginfo_t` structure that describes the signal being delivered.

The `sa_t` field points to an array of pointers to `struct sa_t` which describe the activations involved in this event. The first element of this array, `sa_t[0]`, points to the `sa_t` of the running activation. The next elements of the array (`sa_t[1]` through `sa_t[event+1]`) describe the activations that were preempted, and the thread currently executing on one of the application’s processors is preempted. The upcall is then run in the context of that virtual processor, carrying notifications of both the first thread returning from the kernel and the second thread being preempted. The upcall handler, knowing all of this, can then decide which thread to run next, based on its own processing needs. Typically, this would choose the highest priority thread, which should be one of the two involved in the upcall notification.

The `events` field is a bit mask of events that occurred. The `interrupted` field is a bit mask of events that occurred. The `arg` parameter points to a `siginfo_t` structure that describes the signal being delivered.

The `interrupted` field is a bit mask of events that occurred. The `arg` parameter points to a `siginfo_t` structure that describes the signal being delivered.

The `sa_t` field points to an array of pointers to `struct sa_t` which describe the activations involved in this event. The first element of this array, `sa_t[0]`, points to the `sa_t` of the running activation. The next elements of the array (`sa_t[1]` through `sa_t[event+1]`) describe the activations that were preempted, and the thread currently executing on one of the application’s processors is preempted. The upcall is then run in the context of that virtual processor, carrying notifications of both the first thread returning from the kernel and the second thread being preempted. The upcall handler, knowing all of this, can then decide which thread to run next, based on its own processing needs. Typically, this would choose the highest priority thread, which should be one of the two involved in the upcall notification.

The `events` field is a bit mask of events that occurred. The `interrupted` field is a bit mask of events that occurred. The `arg` parameter points to a `siginfo_t` structure that describes the signal being delivered.

The `interrupted` field is a bit mask of events that occurred. The `arg` parameter points to a `siginfo_t` structure that describes the signal being delivered.

The `sa_t` field points to an array of pointers to `struct sa_t` which describe the activations involved in this event. The first element of this array, `sa_t[0]`, points to the `sa_t` of the running activation. The next elements of the array (`sa_t[1]` through `sa_t[event+1]`) describe the activations that were preempted, and the thread currently executing on one of the application’s processors is preempted. The upcall is then run in the context of that virtual processor, carrying notifications of both the first thread returning from the kernel and the second thread being preempted. The upcall handler, knowing all of this, can then decide which thread to run next, based on its own processing needs. Typically, this would choose the highest priority thread, which should be one of the two involved in the upcall notification.

The `events` field is a bit mask of events that